



CODE DEMO

Capsule Networks



Agenda

- Theory
- Model Architecture
- Layers
- Training
- Testing
- Reconstruction Images



Theory Explained



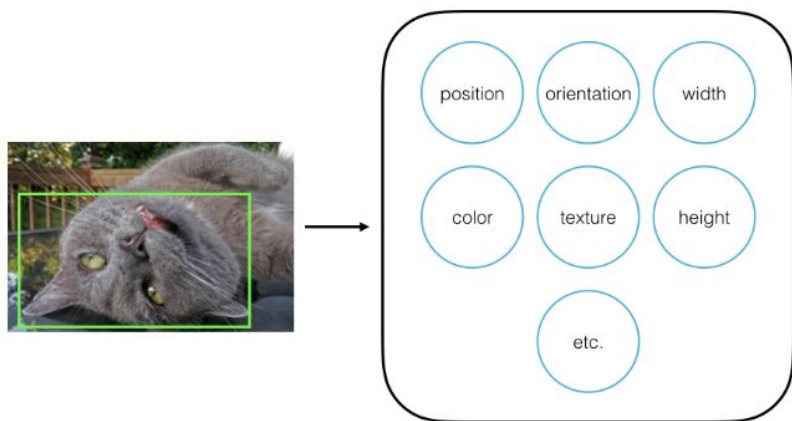
What are Capsules?

Capsules are a small group of neurons that have a few key traits:

Each neuron in a capsule represents various properties of a particular image part; properties like a parts color, width, etc.

Examples:

- Cat: what is the position, orientation texture etc.
- House: What is the angle, what shapes are present etc.



Representing Relationships Between Parts

Dynamic Communication:


- Capsules communicate to determine data flow.
- Enables learning of spatial relationships between parts.

Learning Spatial Relationships:

- Capsule networks learn how parts relate to wholes.
- Facilitates object recognition regardless of orientation.

Advantages Over Vanilla CNN:

- Improved object identification in various orientations.
- Better at recognizing multiple, overlapping objects.
- Enhanced learning from smaller training datasets.



Dataset

MNIST



Dataset

For our demo today, we're using the MNIST dataset.

MNIST consists of 70,000 handwritten digits, each a 28x28 pixel grayscale image.

How to Download and Preprocess the Dataset ?

To get started with MNIST, you can easily download it using TensorFlow or PyTorch with built-in functions.



Model Architecture

Conv Layers

Primary Caps



Digit Caps

Decoder

Encoder Layers


Encoder

The encoder is made of a series of layers that are responsible for taking in as input a 28 by 28 MNIST image and learning to encode it into a 16-dimensional output vector.



First Layer: Convolutional Layer

The first layer in our encoder is a convolutional layer that will learn to extract features, like edges, in a given input image. This convolutional layer will create a stack of 256 filtered images, given one input MNIST image.



Second Layer: Primary Capsules


This layer has 8 primary capsules. Essentially, each capsule is responsible for producing weighted combinations of the features detected in the previous convolutional layer.

Squashing

We define a nonlinear function squash that calculates a certain capsule's normalized, vector output using the following equation.

$$v_j = \frac{\|s_j^2\| s_j}{1 + \|s_j^2\| s_j}$$

v_j is the value we want to calculate, the normalized vector output of a capsule j . And s_j is that capsule's total input; a weighted sum over all the output vectors from the capsules in the layer below capsule j .



Third Layer: Digit Capsules


This layer is composed of 10 "digit" capsules, one for each of our digit classes 0-9. Each capsule takes, as input, a batch of 1152-dimensional vectors produced by our 8 primary capsules, above.

Dynamic Routing

- Facilitates finding the optimal connections between capsules in different layers.
- Enables capsules to communicate and determine the flow of data.
- Ensures that relevant information is routed to appropriate parent capsules.

Coupling Coefficients

- Represent the probability that a child capsule's output should be routed to a parent capsule.
- Initially uncertain, coupling coefficients are determined through dynamic routing.
- Sum of coupling coefficients across all possible parent capsules equals 1.



Third Layer: Digit Capsules

Routing by Agreement


- Iterative process updating coupling coefficients between capsules.
- Determines the flow of information between capsule layers.
- Child capsule outputs a vector indicating part existence and pose.
- Computes prediction vector for each possible parent capsule.
- Agreement between prediction and parent vectors increases coupling coefficients.
- High coupling coefficient enhances the child's influence on the parent.
- Capsules iteratively refine their connections based on agreement.

$$c_{ij} = \frac{e^{b_{ij}}}{\sum_k e^{b_{ik}}}$$

$$\hat{u} = Wu$$

$$a = v \cdot u$$


$$b_{ij} = b_{ij} + a$$



Third Layer: Digit Capsules

Digit Capsules

- DigitCaps layer consists of 10 capsules, each representing a digit class (0-9).
- Each capsule takes 1152-dimensional vectors from the primary capsules as input.
- Output of each capsule is a 16-dimensional vector.



Decoder

Decoder Explained

- The decoder takes 16-dimensional vectors from the DigitCaps layer as input.
- Identifies the "correct" capsule output vector, which corresponds to the digit capsule with the largest vector magnitude.
- Measures the difference between the input image and the decoder reconstruction using Euclidean distance.




Image from [the original Capsule Network paper](#)



Create the Complete Model

- Includes convolutional layers, primary and digit capsule layers, and a decoder.




Image from [the original Capsule Network paper](#)


```
→ CapsuleNetwork(  
    (conv_layer): ConvLayer(  
        (conv): Conv2d(1, 256, kernel_size=(9, 9), stride=(1, 1))  
    )  
    (primary_capsules): PrimaryCaps(  
        (capsules): ModuleList(  
            (0-7): 8 x Conv2d(256, 32, kernel_size=(9, 9), stride=(2, 2))  
        )  
    )  
    (digit_capsules): DigitCaps()  
    (decoder): Decoder(  
        (linear_layers): Sequential(  
            (0): Linear(in_features=160, out_features=512, bias=True)  
            (1): ReLU(inplace=True)  
            (2): Linear(in_features=512, out_features=1024, bias=True)  
            (3): ReLU(inplace=True)  
            (4): Linear(in_features=1024, out_features=784, bias=True)  
            (5): Sigmoid()  
        )  
    )  
)
```

Custom Loss

Margin Loss & Reconstruction Loss


- Custom loss function for training the capsule network.
- Incorporates two key components: margin loss and reconstruction loss.
- Margin loss penalizes deviations from ideal capsule vector magnitudes based on class labels.
- Reconstruction loss measures the dissimilarity between original input images and reconstructed images.
- Combined loss function guides the training process to optimize both classification performance and reconstruction quality.

$$L_c = T_c(\max[0, 0.9 - v_c]) + \lambda(1 - T_c)(\max[0, v_c - 0.1])$$



Loss Function and Optimizer


- Utilize a custom loss function, defined as the **CapsuleLoss**, to optimize the capsule network.
- The CapsuleLoss incorporates both margin loss and reconstruction loss to train the network effectively.
- Implement the Adam optimizer with default parameters for efficient parameter updates during training.



Train the Network

Steps

- Clear gradients, forward pass, calculate loss, backward pass, and optimization steps are performed in each iteration.
- Training loss is recorded and printed intermittently to monitor training progress.
- We are providing insights into how the model learns over time and improves its performance.
- Plot the recorded training loss to visualize how it decreases over the epochs, indicating model improvement.



Test the Trained Network


Testing

- Using the test data to evaluate the performance of the trained model.
- Evaluate classification accuracy and loss metrics on unseen test data.
- Granular analysis of performance on each class.
- Provides insights into how well the model generalizes to new data.

➡ Test Loss: 0.03409957

```
Test Accuracy of Class 0: 99.69% (977/980)
Test Accuracy of Class 1: 99.65% (1131/1135)
Test Accuracy of Class 2: 98.64% (1018/1032)
Test Accuracy of Class 3: 97.92% (989/1010)
Test Accuracy of Class 4: 98.27% (965/982)
Test Accuracy of Class 5: 98.99% (883/892)
Test Accuracy of Class 6: 98.85% (947/958)
Test Accuracy of Class 7: 99.51% (1023/1028)
Test Accuracy of Class 8: 99.08% (965/974)
Test Accuracy of Class 9: 98.32% (992/1009)
```


Overall Test Accuracy: 98.90% (9890/10000)



Display Reconstructions


Image Reconstructions

- Display the original MNIST images and their reconstructions to evaluate decoder performance.
- We use a function to plot one row of original images and another row of their corresponding reconstructions.
- Reconstructions appear nearly identical to the original images, with slight blurring/smoothing at the edges.



Task

Implement Capsule Networks on a New Dataset: You will apply the capsule network architecture to alternative datasets such as Fashion-MNIST or CIFAR-10. This exercise will enhance understanding of the model's adaptability across various image data types and allow exploration of its efficacy in more intricate settings.



The background features a complex pattern of blue lines and arrows. Solid lines intersect at various angles, while dashed lines form loops and curves. Small circles, some open and some filled, are placed at various points along these lines, suggesting a network or a path. The overall aesthetic is technical and modern.

Thankyou